

الجمهورية الجزائرية الديمقراطية الشعبية
وزارة التعليم العالي والبحث العلمي
المدرسة العليا للإعلام الآلي 80 ماي 5491 - سيدي بلعباس

People's Democratic Republic of Algeria
Ministry of Higher Education and Scientific Research
Higher School of Computer Science – May 8, 1945 – Sidi Bel Abbès



Master's Thesis

In partial fulfillment of the requirements for the Master degree in Computer Science

Specialization: Information Systems and Technologies

Topic

AI-Assisted Semantic Data Integration

A Study of Artificial Intelligence Approaches for Enhancing Semantic Web Data Interoperability

Presented by

Hakim ALLEM

Defended in: **10 March 2025**

In front of the jury composed of:

Mr. Mohamed Yacine Kazi Tani

President of the Jury

Mr. Mohamed Bedjaoui (ESI-SBA)

Academic Supervisor

Mrs. Christelle Pierkot (Data Terra)

Professional Supervisor

Mrs. Khayra Benchrif

Examiner

Acknowledgement

This report marks the culmination of five years of study at the **Higher School of Computer Science (ESI-SBA)**, years that have shaped me both technically and personally. I am deeply grateful to the entire ESI community for their support, encouragement, and the high-quality education I have received.

I would like to express my sincere gratitude to my supervisors, **Dr. Mohamed Bedjaoui** and **Mrs. Christelle Pierkot**, for their guidance, availability, and insightful advice throughout this internship. Their mentorship has greatly contributed to the successful completion of this project and to my professional development.

I also extend my thanks to the entire team at the host organization for their warm welcome and the collaborative environment that enriched my experience during the internship.

To my family, thank you for your unwavering love, patience, and support. Your presence has been my strongest source of motivation. I am also grateful to my friends and fellow students who have been by my side through challenges and successes—thank you for the discussions, encouragement, and unforgettable memories.

Finally, I am thankful to everyone who has played a role, directly or indirectly, in this journey. Your contributions, no matter how small, have made this achievement possible.

Abstract

The rapid growth of heterogeneous and distributed data has created a critical need for advanced integration approaches that go beyond traditional syntactic and structural matching. **AI-assisted Semantic Data Integration** (AI-SDI) combines Semantic Web standards with artificial intelligence techniques to automate and enhance the alignment, merging, and enrichment of data from diverse sources.

This report provides a comprehensive study of AI-SDI, examining its theoretical foundations, key enabling technologies, and recent research trends. We explore how AI techniques—such as machine learning, natural language processing, and knowledge graph embeddings—can improve ontology alignment, entity resolution, schema matching, and semantic enrichment. The analysis includes a review of state-of-the-art approaches, their strengths and limitations, and the role of human-in-the-loop strategies for ensuring accuracy and explainability.

The study highlights both the opportunities and challenges of integrating AI into semantic workflows, including scalability, interoperability, and ethical considerations. By synthesizing existing literature and identifying promising research directions, this work aims to provide a reference point for future developments in AI-assisted semantic data integration.

Keywords— Semantic Data Integration, Artificial Intelligence, Ontology Alignment, Knowledge Graphs, Semantic Web, Machine Learning

Résumé

La croissance rapide des données hétérogènes et distribuées crée un besoin crucial de méthodes d'intégration avancées allant au-delà des approches purement syntaxiques ou structurelles. **L'intégration sémantique des données assistée par l'IA** (AI-SDI) combine les standards du Web sémantique avec des techniques d'intelligence artificielle pour automatiser et améliorer l'alignement, la fusion et l'enrichissement de données provenant de sources diverses.

Ce mémoire présente une étude complète d'AI-SDI, en examinant ses fondements théoriques, ses technologies clés et les tendances récentes de la recherche. Nous explorons comment les techniques d'IA—telles que l'apprentissage automatique, le traitement du langage naturel et les plongements de graphes de connaissances—peuvent améliorer l'alignement d'ontologies, la résolution d'entités, la correspondance de schémas et l'enrichissement sémantique. L'analyse inclut un examen des approches de pointe, de leurs avantages et limites, ainsi que du rôle des stratégies *human-in-the-loop* pour garantir la précision et l'explicabilité.

L'étude met en lumière les opportunités et les défis liés à l'intégration de l'IA dans les flux de travail sémantiques, y compris la montée en charge, l'interopérabilité et les considérations éthiques. En synthétisant la littérature existante et en identifiant les perspectives de recherche prometteuses, ce travail constitue une référence pour les futurs développements dans le domaine de l'intégration sémantique des données assistée par l'IA.

Mots-clés— Intégration sémantique des données, Intelligence artificielle, Alignement d'ontologies, Graphes de connaissances, Web sémantique, Apprentissage automatique

Arabic abstract

إن النمو السريع للبيانات الموزعة والمتنوعة المصدر يفرض حاجة ملحة إلى أساليب متقدمة لدمج البيانات تتجاوز المطابقة التركيبية أو النحوية التقليدية. تكامل البيانات الدلالية بمساعدة الذكاء الاصطناعي يجمع بين معايير الويب الدلالي وتقنيات الذكاء الاصطناعي لأتمتة وتحسين عمليات مواءمة البيانات ودمجها وإثرائها من مصادر مختلفة.

يقدم هذا العمل دراسة شاملة حول تكامل البيانات الدلالية بمساعدة الذكاء الاصطناعي، حيث يستعرض الأسس النظرية والتقنيات الرئيسية والاتجاهات الحديثة في هذا المجال. كما يستكشف كيفية مساهمة تقنيات الذكاء الاصطناعي—مثل التعلم الآلي، ومعالجة اللغة الطبيعية، وتضمين الرسوم البيانية المعرفية—في تحسين مواءمة الأنطولوجيات، وحل الكيانات، ومطابقة المخططات، والإثراء الدلالي. وتشمل الدراسة مراجعة لأحدث الأساليب، نقاط قوتها وضعفها، ودور استراتيجيات pool-eht-ni لضمان الدقة وقابلية التفسير.

تسلط الدراسة الضوء على الفرص والتحديات المرتبطة بدمج الذكاء الاصطناعي في سير العمل الدلالي، بما في ذلك قابلية التوسع، وقابلية التشغيل البيئي، والاعتبارات الأخلاقية. ومن خلال تلخيص الأبحاث الحالية وتحديد الاتجاهات الواعدة، يوفر هذا العمل مرجعاً لتطويرات مستقبلية في مجال تكامل البيانات الدلالية بمساعدة الذكاء الاصطناعي.

الكلمات المفتاحية— تكامل البيانات الدلالية، الذكاء الاصطناعي، مواءمة الأنطولوجيات، الرسوم البيانية المعرفية، الويب الدلالي، التعلم الآلي

1	Introduction	10
1.1	Background and Motivation	10
1.2	Problem Statement	10
1.3	Objectives	10
1.4	Methodology	11
1.5	Report Structure	11
2	Background Concepts	12
2.1	The Semantic Web	12
2.2	Ontologies	13
2.3	Fundamentals of Data Integration	14
2.4	Challenges in Semantic Data Integration	15
2.5	Summary	15
3	Classical Approaches to Semantic Data Integration	16
3.1	Ontology-Based Integration Architectures	16
3.2	Schema and Ontology Matching	17
3.3	Mapping Generation and Maintenance	17
3.4	Query Processing in Semantic Integration	18
3.5	Limitations of Classical Approaches	19
3.6	Summary	19

4	AI-Assisted Semantic Data Integration	20
4.1	Motivation for AI in SDI	20
4.2	Machine Learning for Schema and Ontology Matching	21
4.3	Natural Language Processing for Semantic Enrichment	21
4.4	Knowledge Graph Embeddings	22
4.5	Active Learning and Human-in-the-Loop Approaches	22
4.6	AI for Mapping Maintenance and Evolution	23
4.7	Challenges and Limitations	23
4.8	Summary	24
5	Case Studies and Applications of AI in Semantic Data Integration	25
5.1	Biomedical Data Integration	25
5.2	Geosciences and Environmental Data	26
5.3	Cultural Heritage and Digital Humanities	26
5.4	Enterprise Knowledge Integration	27
5.5	Lessons Learned from Case Studies	27
5.6	Summary	27
6	Challenges, Limitations, and Future Directions	29
6.1	Data Quality and Heterogeneity	29
6.2	Explainability and Trust in AI Systems	29
6.3	Scalability and Performance	30
6.4	Domain Adaptation and Transferability	30
6.5	Ethical and Legal Considerations	30
6.6	Future Directions	31
6.6.1	Improved Explainability	31
6.6.2	Incremental and Lifelong Learning	31
6.6.3	Human-in-the-Loop Integration	31
6.6.4	Standardization and Interoperability	31
6.6.5	Ethics-Aware Integration	32
6.7	Summary	32

LIST OF FIGURES

2.1	The W3C “Semantic Web layer cake” illustrating how identifiers (URI/IRI), RDF/RDFS, OWL, SPARQL, rules and logic stack to enable machine-readable knowledge on the Web.	13
3.1	Global-as-View (GAV) paradigm — the global schema is expressed directly in terms of source queries.	18
3.2	Local-as-View (LAV) paradigm — source schemas are expressed in terms of the global schema.	18
3.3	GLAV (hybrid) paradigm — mappings are expressed as conjunctive queries between global and source schemas.	19
4.1	An AI-assisted semantic mapping pipeline: from heterogeneous inputs and preprocessing to feature extraction, similarity computation, model prediction, and expert validation.	21
4.2	Entity Resolution (ER) and Entity Linking (EL) pipeline for enriching ontologies with external knowledge and resolving equivalent entities.	22
4.3	Knowledge graph embedding approaches: projecting entities and relations into vector spaces for similarity computation and link prediction.	23

	LIST OF TABLES
--	----------------

1.1 Background and Motivation

Semantic Data Integration (SDI) aims to combine heterogeneous data sources into a unified, semantically consistent representation [1]. With the rapid growth of data in various domains, traditional rule-based integration methods face challenges in scalability, adaptability, and handling unstructured or semi-structured data. Recent advances in Artificial Intelligence (AI), particularly in Machine Learning (ML) and Natural Language Processing (NLP), have opened new possibilities for improving semantic mapping, entity resolution, and ontology alignment [knoblock2012exploiting, rahm2001survey].

1.2 Problem Statement

Despite progress in both semantic web technologies and AI techniques, combining them effectively for large-scale data integration remains complex. Challenges include heterogeneous schemas, inconsistent ontologies, domain-specific vocabularies, and the need for real-time integration.

1.3 Objectives

The main objectives of this report are:

- To review the state-of-the-art in Semantic Data Integration.
- To examine AI-assisted techniques for ontology alignment, schema mapping, and entity

resolution.

- To identify potential applications and research gaps in AI-assisted SDI.

1.4 Methodology

This work is based on a literature review of academic articles, technical reports, and standards related to semantic web, data integration, and AI-assisted methods. No original implementation or experimentation is conducted.

1.5 Report Structure

The remainder of this report is organized as follows:

- Chapter 2 covers background concepts in semantic web, ontologies, and data integration.
- Chapter ?? presents traditional semantic data integration methods.
- Chapter ?? explores AI-assisted approaches.
- Chapter ?? reviews applications and case studies.
- Chapter ?? discusses challenges and future research directions.
- Chapter 7 concludes the report.

CHAPTER 2

BACKGROUND CONCEPTS

This chapter introduces the theoretical foundations of AI-Assisted Semantic Data Integration (SDI). It first outlines the vision and principles of the Semantic Web, then explores the role of ontologies as formalized conceptual models, before discussing the fundamentals of data integration and the limitations of current semantic approaches. This background sets the stage for understanding how artificial intelligence can enhance semantic interoperability.

2.1 The Semantic Web

The Semantic Web, introduced by Berners-Lee, Hendler, and Lassila [2], extends the traditional Web by enabling information to be given explicit meaning, allowing both humans and machines to understand and process it. The central idea is to transform the Web from a collection of interlinked documents into a *web of data* where resources are described using formal semantics. By relying on shared standards, the Semantic Web facilitates interoperability between heterogeneous systems and supports advanced services such as intelligent search, automated reasoning, and data integration across distributed sources.

One of the key architectural views of the Semantic Web is represented by the so-called “layer cake” of W3C technologies. At the base are URIs and Unicode, which ensure unique identification and universal encoding of resources. Above them, XML and RDF provide syntax and a graph-based data model, respectively. The vocabulary layer is supported by RDF Schema (RDFS) and the Web Ontology Language (OWL), which enable the definition of domain-specific concepts and relationships. On top, SPARQL acts as the standard query language for RDF datasets, while logic and proof layers allow more complex reasoning capabilities. Together, these technologies form a coherent stack that supports machine-readable knowledge representation on the Web.




Figure 2.1: The W3C “Semantic Web layer cake” illustrating how identifiers (URI/IRI), RDF/RDFS, OWL, SPARQL, rules and logic stack to enable machine-readable knowledge on the Web.

2.2 Ontologies

Ontologies are the cornerstone of the Semantic Web and, by extension, of semantic data integration. Gruber [3] defines an ontology as a “formal, explicit specification of a shared conceptualization.” In practice, an ontology provides a controlled vocabulary of terms, together with formal definitions of their meanings and the relationships between them. This enables different systems and communities to share a common understanding of a domain, even if their underlying data structures differ.

An ontology typically contains classes (representing concepts), properties (defining attributes and relationships), individuals (specific instances), and axioms (logical statements

that constrain possible interpretations). These elements work together to create a semantic structure that supports reasoning and knowledge inference. For instance, in an environmental science ontology, one might define a class **River** as a subclass of **WaterBody**, which allows systems to infer that every river is also a water body.

Several formal languages exist for defining ontologies. The Web Ontology Language (OWL) [mcguinness2004owl] offers varying levels of expressiveness, from OWL Lite for simpler taxonomies to OWL DL for more expressive yet computationally decidable models. SKOS (Simple Knowledge Organization System) [miles2009skos] provides a lightweight model for representing thesauri, classification schemes, and taxonomies. Unlike OWL, SKOS is designed for hierarchical and associative relationships without the complexity of full logical reasoning.

2.3 Fundamentals of Data Integration

Data integration is the process of combining information from multiple, often heterogeneous, sources into a unified and coherent view [1]. In the context of semantic data integration, this process relies on ontologies to serve as the common conceptual model that mediates between differing data schemas. The challenge lies in bridging differences not only in syntax but also in structure and meaning.

Different architectural paradigms exist for semantic data integration. In the Global-as-View (GAV) approach, the global schema is expressed directly in terms of the source schemas. This makes querying straightforward but requires schema changes whenever sources evolve. Conversely, the Local-as-View (LAV) paradigm defines each source schema in terms of the global schema, making it easier to add or modify sources at the cost of more complex query processing. More recently, Ontology-Based Data Access (OBDA) [4] has emerged, where an ontology acts as the global schema, and declarative mappings connect it to the data sources. This allows queries to be expressed in terms of high-level concepts rather than source-specific details.

The process of semantic data integration typically involves schema matching (finding correspondences between elements of different schemas) [rahm2001survey], ontology alignment (identifying semantic correspondences between ontologies) [5], and entity resolution (determining when different records refer to the same real-world entity) [6]. Each of these tasks requires both structural analysis and semantic interpretation, often involving human expertise.

2.4 Challenges in Semantic Data Integration

While the Semantic Web provides a strong theoretical foundation for interoperability, traditional semantic data integration approaches face several persistent challenges. Data heterogeneity is perhaps the most fundamental: sources may differ in structure (relational tables vs. RDF graphs), syntax (different encodings or formats), and semantics (different interpretations of the same term). Scalability is another concern, as reasoning and matching algorithms can become computationally expensive when dealing with very large datasets or complex ontologies.

Dynamic environments add further complexity, as data sources and ontologies evolve over time. Changes in terminology, schema structure, or conceptual scope require continuous maintenance of mappings and alignments. Moreover, many semantic integration tasks still depend heavily on manual effort by domain experts, particularly when dealing with ambiguous mappings or incomplete ontological coverage. These limitations motivate the exploration of artificial intelligence techniques to automate, optimize, and scale the integration process.

2.5 Summary

This chapter has presented the core concepts underlying AI-Assisted Semantic Data Integration. The Semantic Web offers a standardized technological framework for representing and linking data, while ontologies provide the formal semantics needed for interoperability. Data integration principles and architectures, along with their associated challenges, set the context for why AI techniques are increasingly being applied to semantic integration problems. The next chapter will explore classical semantic data integration methods in more detail, serving as a basis for understanding how AI can enhance or replace them.

CHAPTER 3

CLASSICAL APPROACHES TO SEMANTIC DATA INTEGRATION

Before exploring how artificial intelligence can enhance semantic data integration (SDI), it is important to understand the classical approaches that form its foundation. These methods emerged in the early development of the Semantic Web and remain widely used in both academic research and practical applications. This chapter reviews ontology-based data integration architectures, schema and ontology matching techniques, and the associated mapping and query processing strategies. It also discusses their strengths and limitations, which later motivate AI-driven enhancements.

3.1 Ontology-Based Integration Architectures

Ontology-Based Data Integration (OBDI) is the dominant paradigm in semantic interoperability. In this approach, an ontology acts as a shared conceptual model that mediates between heterogeneous data sources. Each source is connected to the ontology through declarative mappings, allowing queries to be expressed at the conceptual level rather than in terms of specific database schemas [7, 4].

Two main architectural patterns have historically guided OBDI: *single-ontology* and *multiple-ontology* approaches. In the single-ontology approach, all sources map to one common ontology that captures the domain semantics. This promotes semantic consistency but requires significant effort to design and maintain the global ontology, especially in domains with diverse conceptualizations. In contrast, the multiple-ontology approach allows each source to have its own ontology, with mappings established between them. This reduces the burden of creating a single unifying model but increases the complexity of managing inter-ontology alignments.

A third model, known as the *hybrid approach*, combines these strategies by linking local ontologies to a shared upper ontology while preserving source-specific detail. Hybrid models are particularly useful in large-scale integration projects where full semantic unification is impractical, such as cross-disciplinary scientific data portals.

3.2 Schema and Ontology Matching

One of the core tasks in semantic data integration is matching—establishing correspondences between entities in different schemas or ontologies [rahm2001survey, 5]. Matching can operate at different levels: between database schemas (schema matching) or between formal ontologies (ontology alignment). The goal is to determine when two entities—such as classes, properties, or instances—are semantically equivalent, subsumed, or otherwise related.

Classical matching techniques can be broadly categorized into *lexical*, *structural*, and *instance-based* approaches. Lexical matching relies on string similarity measures, dictionaries, and linguistic resources to compare entity names and labels. Structural matching exploits the relationships and hierarchies within schemas or ontologies, reasoning that elements in similar positions or with similar neighbors are likely to correspond. Instance-based matching uses data values associated with entities to identify correspondences, for example, by detecting overlapping sets of instances between two classes.

In practice, integration systems often combine multiple matching techniques in a hybrid strategy, leveraging the strengths of each. For instance, lexical similarity might suggest candidate correspondences, which are then validated through structural consistency checks. Some systems also incorporate domain-specific thesauri or manually curated mappings to improve precision.

3.3 Mapping Generation and Maintenance

Once correspondences between entities have been established, they must be expressed as formal mappings that can be used for query rewriting or data transformation. In the Semantic Web context, mappings are often expressed in standards such as R2RML (for relational-to-RDF mappings) [8] or alignment formats defined by the Ontology Alignment Evaluation Initiative (OAEI) [9].

Mapping generation is traditionally a semi-automatic process: algorithms produce candidate mappings, which are then reviewed and validated by domain experts. This manual validation step remains essential in critical applications, where incorrect mappings can lead to erroneous data integration or misleading results. Maintenance is an ongoing challenge, as changes to source schemas or ontologies may render existing mappings obsolete. Versioning

mechanisms and incremental matching techniques have been proposed to reduce the cost of maintaining mappings in dynamic environments.

3.4 Query Processing in Semantic Integration

In ontology-based systems, queries are typically expressed in SPARQL against the global ontology. Query processing involves translating these high-level queries into source-specific queries, a process known as query rewriting [10]. In the GAV (Global-as-View) paradigm, rewriting is relatively straightforward, as each global concept has a direct definition in terms of the sources. In LAV (Local-as-View), rewriting requires reasoning to identify which sources can contribute to answering a query, often resulting in complex query plans.




Figure 3.1: Global-as-View (GAV) paradigm — the global schema is expressed directly in terms of source queries.



Figure 3.2: Local-as-View (LAV) paradigm — source schemas are expressed in terms of the global schema.

Performance optimization is a major concern in semantic query processing, particularly when integrating large or distributed datasets. Techniques such as caching intermediate results, indexing semantic concepts, and pushing filters to source queries are commonly employed. In some cases, materialized integration—where integrated data is precomputed and stored—can offer performance benefits at the cost of freshness.




Figure 3.3: GLAV (hybrid) paradigm — mappings are expressed as conjunctive queries between global and source schemas.

3.5 Limitations of Classical Approaches

Despite their maturity, classical semantic data integration approaches face persistent limitations that motivate the adoption of AI-assisted methods. The reliance on manual mapping validation remains a bottleneck, especially when integrating large numbers of sources. Matching algorithms, while effective in many cases, can struggle with domain-specific terminology, evolving vocabularies, or sparse data. Scalability issues emerge as ontologies grow in size and complexity, and maintaining consistency in dynamic environments is costly.

Furthermore, classical systems often lack adaptability: once mappings and integration rules are established, they do not automatically improve or adjust unless explicitly updated. In rapidly evolving domains, this can lead to outdated or suboptimal integration. These challenges create opportunities for AI techniques—such as machine learning, natural language processing, and knowledge graph embeddings—to enhance automation, accuracy, and adaptability in semantic data integration.

3.6 Summary

This chapter has reviewed the principal components of classical semantic data integration: ontology-based architectures, matching techniques, mapping generation, and query processing. While these methods form a robust foundation for interoperability, their limitations in scalability, adaptability, and automation suggest the need for complementary AI-driven approaches. The next chapter will examine how artificial intelligence techniques can be integrated into the semantic data integration pipeline to address these challenges and extend its capabilities.

CHAPTER 4

AI-ASSISTED SEMANTIC DATA INTEGRATION

Artificial Intelligence (AI) has emerged as a transformative force in semantic data integration (SDI), offering solutions to many of the limitations inherent in classical approaches. While traditional methods rely heavily on manually engineered mappings and static rules, AI introduces adaptive, data-driven techniques capable of learning correspondences, improving over time, and handling scale and complexity in ways that were previously impractical. This chapter examines how AI techniques—particularly machine learning, natural language processing, and knowledge graph embeddings—are applied across the SDI pipeline, from schema and ontology matching to mapping maintenance and semantic enrichment.

4.1 Motivation for AI in SDI

The motivation for integrating AI into SDI stems from three main factors: automation, scalability, and adaptability. Manual mapping creation and validation, while precise, is time-consuming and prone to human error, especially in domains with thousands of heterogeneous data sources. AI-based models can automate much of this process by learning patterns from existing mappings or from textual and structural features of schemas and ontologies. Scalability is another driver: AI algorithms, especially those leveraging distributed computation, can handle large-scale integration scenarios involving millions of entities. Finally, adaptability is crucial in dynamic environments where vocabularies evolve rapidly; AI models can be retrained or fine-tuned to reflect new semantic relationships without rewriting the integration logic from scratch.

4.2 Machine Learning for Schema and Ontology Matching

One of the earliest applications of AI in SDI is the use of machine learning to improve schema and ontology matching. Instead of relying solely on string similarity or structural heuristics, supervised and semi-supervised models learn from labeled correspondences to predict new matches [11, 12]. Features can include lexical similarity scores, graph neighborhood statistics, ontology depth, and instance data overlap. Ensemble approaches combine multiple similarity metrics into a single predictive model, often outperforming individual heuristics.

Recent work has applied deep learning to this task, using representation learning to capture semantic similarities between entities. For example, neural models can embed labels and definitions into vector spaces using word embeddings such as Word2Vec [13], GloVe [14], or contextual embeddings from models like BERT [15]. These embeddings capture rich semantic relationships, enabling matches even when labels differ significantly in wording.




Figure 4.1: An AI-assisted semantic mapping pipeline: from heterogeneous inputs and pre-processing to feature extraction, similarity computation, model prediction, and expert validation.

4.3 Natural Language Processing for Semantic Enrichment

Natural Language Processing (NLP) techniques are particularly valuable for extracting semantic cues from textual descriptions associated with schemas and ontologies. Many ontologies contain definitions, scope notes, and usage examples that can be processed with NLP pipelines to derive domain-specific terminology and detect semantic similarities.

Entity linking methods can align terms in descriptions to external knowledge bases such as DBpedia [16] or Wikidata [17], thereby enriching ontologies with additional context. Furthermore, relation extraction techniques can identify implicit semantic relationships between entities, expanding the integration beyond explicitly modeled connections.

Transformer-based language models have significantly advanced NLP for SDI, as they can model complex dependencies in text and adapt to specialized domains through fine-tuning.

In semantic integration, such models can be trained to perform tasks like cross-ontology concept alignment or classification of mappings into equivalence, subsumption, or relatedness categories.



Figure 4.2: Entity Resolution (ER) and Entity Linking (EL) pipeline for enriching ontologies with external knowledge and resolving equivalent entities.

4.4 Knowledge Graph Embeddings

Knowledge graph embedding techniques map entities and relations in a graph to low-dimensional vector spaces while preserving structural and semantic properties [18]. These embeddings enable efficient similarity computation, clustering, and reasoning over large-scale integrated datasets. Models such as TransE [19], DistMult [20], and ComplEx [21] have been widely adopted for link prediction tasks, which can be adapted to suggest potential mappings between ontologies.

By learning embeddings jointly over multiple sources, AI systems can capture both intra-source and inter-source semantic patterns, facilitating more robust integration. Embedding-based matching is particularly effective in scenarios where schema or ontology labels are sparse or non-informative, as the model relies on relational structure rather than textual cues alone.

4.5 Active Learning and Human-in-the-Loop Approaches

While AI can automate many aspects of SDI, human expertise remains essential for ensuring semantic correctness. Active learning strategies integrate human feedback into the AI model training process, selecting the most uncertain or potentially informative mapping candidates for expert review [22]. This approach reduces the annotation burden by focusing expert effort where it has the greatest impact.

Human-in-the-loop systems for SDI not only validate mappings but also provide iterative corrections that can be used to retrain AI models. This continuous feedback loop enables the integration system to adapt to evolving vocabularies and domain-specific nuances, achieving a balance between automation and expert oversight.

Knowledge Graph Embeddings (2D projection)




Figure 4.3: Knowledge graph embedding approaches: projecting entities and relations into vector spaces for similarity computation and link prediction.

4.6 AI for Mapping Maintenance and Evolution

Mappings are not static; they must evolve as schemas and ontologies change. AI can assist in mapping maintenance by detecting changes in the source or target models and predicting which mappings are likely to be affected. Incremental learning techniques allow models to update their predictions without retraining from scratch, enabling near-real-time adaptation in dynamic environments.

Change detection can be supported by ontology versioning systems, where AI analyzes differences between ontology versions and suggests mapping updates accordingly. This capability is particularly important in rapidly evolving domains such as biomedical research, where controlled vocabularies are updated frequently.

4.7 Challenges and Limitations

Despite its promise, AI-assisted SDI faces significant challenges. Training supervised models requires labeled datasets, which can be expensive to obtain at scale. Domain adaptation remains a concern: models trained on one domain may not perform well in another without fine-tuning. Interpretability is also a key issue; AI models, particularly deep learning systems, can act as "black boxes," making it difficult to justify specific mapping decisions. Furthermore, computational requirements for training and inference can be high, especially

for transformer-based models or large-scale knowledge graph embeddings.

Another limitation is that AI-based approaches do not eliminate the need for semantic grounding in domain ontologies; rather, they complement it. The best results are often achieved by combining AI with established Semantic Web principles, ensuring that learned mappings are consistent with formal domain models and reasoning rules.

4.8 Summary

AI-assisted semantic data integration builds on classical approaches by introducing automation, adaptability, and scalability into the integration pipeline. Techniques from machine learning, NLP, and knowledge graph embeddings can enhance schema matching, semantic enrichment, mapping maintenance, and query processing. However, these benefits come with challenges in terms of data requirements, interpretability, and computational cost. In the next chapter, we will examine concrete case studies and systems that successfully integrate AI into semantic data integration workflows, highlighting their architectures, methodologies, and results.

CHAPTER 5

CASE STUDIES AND APPLICATIONS OF AI IN SEMANTIC DATA INTEGRATION

While the previous chapter discussed the theoretical foundations and methodologies of AI-assisted semantic data integration (SDI), this chapter focuses on real-world systems and projects where these techniques have been successfully implemented. These case studies illustrate how AI can be combined with Semantic Web technologies to address integration challenges in domains such as biomedical research, geosciences, cultural heritage, and enterprise data management.

5.1 Biomedical Data Integration

The biomedical domain has long faced the challenge of integrating heterogeneous datasets originating from diverse sources such as clinical trials, genomic studies, and biomedical literature. Ontologies such as SNOMED CT, the Gene Ontology, and MeSH provide a strong semantic foundation, yet the volume and complexity of biomedical data necessitate AI support.

One notable example is the **NCBO BioPortal** platform [23], which hosts hundreds of biomedical ontologies. AI-based annotators process unstructured biomedical texts to automatically link terms to ontology concepts, supporting semantic enrichment and discovery. Machine learning models are used to improve the accuracy of these annotations by learning from curator feedback, thereby reducing the need for extensive manual tagging.

In the **OpenPHACTS** project [24], AI-driven entity recognition and ontology alignment techniques were employed to integrate pharmacological data from multiple sources. The system uses text mining to extract drug-target-disease relationships and maps them to

linked data formats, enabling cross-dataset queries. This approach demonstrates how AI and linked data principles can be combined to accelerate drug discovery.

5.2 Geosciences and Environmental Data

In environmental and Earth science research, datasets are often collected from diverse sensors, observation systems, and modeling platforms, leading to significant heterogeneity in formats and semantics. The **Earth Science Knowledge Graph (ESKG)** initiative [25] leverages AI techniques to automatically integrate metadata from multiple sources such as NASA and NOAA repositories. Using natural language processing and ontology alignment, ESKG constructs an interlinked knowledge graph that supports interdisciplinary queries across domains like oceanography, climate, and geology.

Similarly, the **GeoKnow** project [26] applies AI-assisted link discovery to integrate geospatial data on the Web. Geospatial-specific similarity measures are combined with machine learning classifiers to discover links between entities in different datasets, improving search and visualization capabilities. The resulting knowledge graphs facilitate semantic interoperability between environmental monitoring systems and spatial data infrastructures.

5.3 Cultural Heritage and Digital Humanities

The cultural heritage sector faces the challenge of integrating heterogeneous catalogues, museum databases, and archival resources. AI-assisted SDI approaches have been used to link artefact descriptions, bibliographic records, and multimedia content to shared reference vocabularies such as CIDOC CRM.

The **Europeana** platform [27] integrates metadata from thousands of cultural institutions across Europe. AI-powered entity recognition and multilingual semantic alignment techniques are used to unify records that may be described in different languages and terminologies. By embedding descriptions in a shared vector space, the system can match semantically equivalent items despite lexical and linguistic variations.

Another example is the **Pelagios Commons** initiative [28], which applies AI-assisted named entity recognition to link historical place references in texts and maps to authoritative geospatial datasets. This approach has enabled researchers to explore historical sources through interactive linked data maps.

5.4 Enterprise Knowledge Integration

In the corporate sector, semantic data integration is critical for unifying information across business units, applications, and partners. AI-assisted methods are increasingly used to automate the mapping of enterprise schemas, detect duplicates, and enrich data with external sources.

The ****Google Knowledge Graph**** [29] is a large-scale example of AI-driven integration, combining data from multiple sources such as Freebase, Wikipedia, and proprietary datasets. Entity linking and knowledge graph embeddings allow the system to provide contextual answers in search results. Although not fully open, its architecture demonstrates how AI can scale semantic integration to billions of entities.

In manufacturing and supply chain management, AI-assisted SDI supports interoperability between product catalogues, standards, and sensor data. For instance, semantic matching algorithms can align parts lists from different suppliers, while AI-based anomaly detection ensures data consistency across integrated systems.

5.5 Lessons Learned from Case Studies

Across domains, several common themes emerge from the application of AI-assisted SDI:

- **Semantic grounding is essential:** Even in AI-driven systems, ontologies remain the backbone for ensuring integration consistency.
- **Hybrid methods are most effective:** Combining AI with rule-based approaches often yields better accuracy than either method alone.
- **Domain adaptation is key:** Models trained on one domain must often be fine-tuned to perform well in another, highlighting the importance of domain-specific corpora.
- **Human expertise remains vital:** Active learning and curator feedback loops are crucial for refining AI predictions.

5.6 Summary

These case studies illustrate the versatility of AI-assisted semantic data integration across different domains. From biomedical research to cultural heritage, AI methods such as NLP, knowledge graph embeddings, and machine learning-based ontology matching have enhanced

the automation, scalability, and adaptability of integration processes. While challenges remain—particularly in terms of interpretability, domain adaptation, and data quality—AI-assisted SDI has demonstrated tangible benefits in real-world systems, paving the way for further research and broader adoption.

CHAPTER 6

CHALLENGES, LIMITATIONS, AND FUTURE DIRECTIONS

While AI-assisted semantic data integration (SDI) has demonstrated significant potential, its deployment in real-world scenarios faces a number of technical, methodological, and organizational challenges. This chapter critically examines these limitations and outlines promising avenues for future research.

6.1 Data Quality and Heterogeneity

A fundamental challenge in SDI lies in the quality and heterogeneity of input data. Even when AI models are used to align and integrate datasets, issues such as missing values, inconsistent formats, and ambiguous entity references can significantly reduce integration accuracy [30]. In semantic contexts, low-quality or poorly maintained ontologies exacerbate these problems, as the mappings generated by AI rely on the correctness and completeness of reference vocabularies.

Moreover, semantic resources often vary in granularity and scope, making alignment inherently difficult. For instance, a general-purpose ontology may represent concepts at a broad level, while a domain-specific ontology may include highly detailed subcategories. AI techniques such as ontology embeddings or graph neural networks can partially mitigate this issue, but the risk of misalignment remains, particularly in domains with evolving terminologies.

6.2 Explainability and Trust in AI Systems

One of the most frequently cited limitations of AI-assisted SDI is the lack of transparency in model decision-making. Many state-of-the-art integration systems employ deep learning methods, which—despite high accuracy—are often considered “black boxes” [31]. In con-

texts such as biomedical research or regulatory compliance, stakeholders require explainable mappings to ensure traceability and accountability.

Current approaches to explainable AI (XAI) in semantic integration include attention visualization for transformer-based models, rule extraction from neural systems, and hybrid models that combine machine learning with symbolic reasoning [32]. However, these methods remain underexplored in large-scale SDI deployments, and balancing interpretability with predictive performance continues to be a key research challenge.

6.3 Scalability and Performance

Semantic data integration often operates over massive datasets and ontologies containing millions of entities and relationships. While AI methods such as approximate nearest neighbor search and distributed knowledge graph embeddings can improve performance [18], scalability remains a bottleneck, especially when integration tasks require complex reasoning over multiple ontologies.

Cloud-based and distributed processing frameworks, such as Apache Spark or GPU-accelerated graph databases, are increasingly being adopted to address scalability. Nevertheless, ensuring that these systems maintain semantic correctness while handling high-throughput integration remains a non-trivial challenge.

6.4 Domain Adaptation and Transferability

AI models trained for semantic integration in one domain often fail to generalize to others without significant adaptation. This is particularly problematic for low-resource domains, where there is limited training data or curated ontologies. Transfer learning and few-shot learning approaches [33, 34] have shown promise in reducing retraining costs, but they introduce additional complexity in aligning semantic spaces across domains.

A related issue is the dynamic nature of certain knowledge domains. Ontologies and vocabularies are not static—they evolve to incorporate new terms, relations, and classifications. This evolution can cause “concept drift” in AI models, requiring continuous retraining or incremental learning strategies.

6.5 Ethical and Legal Considerations

As AI-assisted SDI systems integrate data from multiple sources, they often encounter sensitive or proprietary information. This raises concerns about privacy, consent, and intellectual

property rights. For example, integrating patient health records in biomedical contexts requires strict compliance with regulations such as GDPR in the EU or HIPAA in the US [35].

Bias in training data can also propagate into integrated systems, leading to unequal representation of certain concepts or entities. Developing bias detection and mitigation strategies is therefore critical to ensuring that semantic integration systems remain fair and non-discriminatory.

6.6 Future Directions

Future research on AI-assisted semantic data integration should focus on the following priorities:

6.6.1 Improved Explainability

Developing transparent AI methods that can provide justifications for each integration decision is essential. This could include integrating provenance tracking directly into AI models, so that each mapping can be traced back to its supporting evidence.

6.6.2 Incremental and Lifelong Learning

Given the dynamic nature of many domains, there is a need for AI models capable of continuous adaptation without complete retraining. Lifelong learning approaches could enable systems to incorporate new concepts and relations into their integration pipelines in real time.

6.6.3 Human-in-the-Loop Integration

Hybrid workflows that involve human experts in the integration process can improve both accuracy and trust. AI systems can propose candidate mappings, which are then reviewed and validated by domain specialists. Active learning strategies can make this process more efficient by prioritizing cases where AI confidence is low.

6.6.4 Standardization and Interoperability

The development of open benchmarks, evaluation datasets, and standardized APIs for AI-assisted SDI could accelerate research and facilitate the comparison of different approaches.

Interoperability between AI models and Semantic Web standards such as RDF, OWL, and SHACL should remain a core design principle.

6.6.5 Ethics-Aware Integration

Embedding ethical considerations into the design of AI-assisted SDI systems—from bias detection to privacy preservation—will be essential for ensuring societal trust and regulatory compliance.

6.7 Summary

AI-assisted semantic data integration holds the potential to revolutionize how heterogeneous datasets are linked and understood. However, realizing this potential requires addressing persistent challenges related to data quality, explainability, scalability, domain adaptation, and ethical concerns. By combining advances in AI, Semantic Web standards, and human-centered design, future systems can achieve higher accuracy, transparency, and societal impact.

The emergence of Artificial Intelligence (AI) as a driver for semantic data integration (SDI) marks a significant milestone in the evolution of data interoperability. By leveraging advanced machine learning algorithms, knowledge graph embeddings, and natural language processing techniques, AI-assisted SDI enables the automated alignment, enrichment, and integration of heterogeneous datasets in ways that were previously impractical or prohibitively expensive.

Throughout this report, we have explored the foundations of semantic data integration, the role of Semantic Web standards such as RDF, OWL, and SPARQL, and the contributions of AI techniques to enhancing integration accuracy, scalability, and adaptability. We examined the state of the art in AI-assisted SDI, including ontology alignment methods, entity resolution strategies, and hybrid symbolic–statistical approaches. The discussion highlighted how AI can support tasks ranging from schema matching to ontology population, and from semantic search to automated reasoning.

At the same time, we have acknowledged that the integration of AI into semantic workflows is not without challenges. Issues related to data quality, model explainability, computational scalability, and domain adaptation remain open research problems. Ethical and legal considerations, including privacy, intellectual property rights, and bias mitigation, also demand sustained attention if AI-assisted SDI is to be trusted and widely adopted.

The analysis in this work underscores that AI-assisted semantic data integration should be viewed not as a fully autonomous process, but rather as part of a *human-in-the-loop* ecosystem. By combining the precision and contextual understanding of domain experts with the pattern recognition and scalability of AI systems, organizations can build integration pipelines that are both accurate and interpretable.

Looking ahead, the field presents several promising directions. Advances in explainable AI, incremental learning, and ethics-aware design could address many of the current limitations.

The creation of standardized benchmarks, shared ontologies, and open-source integration frameworks would further accelerate innovation and adoption. Moreover, the integration of AI-assisted SDI with emerging technologies such as federated learning and privacy-preserving computation could enable secure and collaborative data integration across institutional and national boundaries.

In conclusion, AI-assisted semantic data integration represents a transformative opportunity for data-driven science, industry, and governance. Its success will depend not only on technological breakthroughs, but also on a commitment to transparency, interoperability, and ethical responsibility. With the right balance of innovation and oversight, AI-assisted SDI can become a cornerstone of the next generation of knowledge infrastructures, enabling seamless, meaningful, and responsible use of the world's data resources.

Bibliography

BIBLIOGRAPHY

- [1] Maurizio Lenzerini. “Data integration: A theoretical perspective.” In: *Proceedings of the twenty-first ACM SIGMOD-SIGACT-SIGART symposium on Principles of database systems* (2002), pp. 233–246.
- [2] Tim Berners-Lee, James Hendler, and Ora Lassila. “The Semantic Web.” In: *Scientific American* 284.5 (2001), pp. 34–43.
- [3] Thomas R Gruber. “A translation approach to portable ontology specifications.” In: *Knowledge Acquisition*. Vol. 5. 2. 1993, pp. 199–220.
- [4] Antonella Poggi et al. “Linking data to ontologies.” In: *Journal on data semantics X* (2008), pp. 133–173.
- [5] Jérôme Euzenat and Pavel Shvaiko. “Ontology matching.” In: *Springer* (2007).
- [6] Peter Christen. *Data Matching: Concepts and Techniques for Record Linkage, Entity Resolution, and Duplicate Detection*. Springer, 2012.
- [7] Holger Wache et al. “Ontology-based integration of information—a survey of existing approaches.” In: *IJCAI-01 Workshop: Ontologies and Information Sharing 2001* (2001), pp. 108–117.
- [8] Souripriya Das, Seema Sundara, and Richard Cyganiak. “R2RML: RDB to RDF mapping language.” In: *W3C recommendation 27* (2012).
- [9] Jérôme Euzenat et al. “Ontology alignment evaluation initiative: six years of experience.” In: *Journal on data semantics*. Springer, 2013, pp. 158–192.
- [10] Diego Calvanese, Giuseppe De Giacomo, Maurizio Lenzerini, et al. “Ontology-based data access and integration.” In: *Proc. of the 34th international conference on current trends in theory and practice of computer science (SOFSEM)*. 2007, pp. 58–69.
- [11] Natalya F Noy and Mark A Musen. “Matching ontologies for context.” In: *Journal of the American Medical Informatics Association* 11.2 (2004), pp. 114–124.

- [12] Thanh Nguyen and Ryutaro Ichise. “A survey of techniques for ontology mapping.” In: *International Journal of Computer Science and Information Technology* 3.3 (2011), pp. 139–160.
- [13] Tomas Mikolov et al. “Distributed representations of words and phrases and their compositionality.” In: *Advances in Neural Information Processing Systems*. Vol. 26. 2013.
- [14] Jeffrey Pennington, Richard Socher, and Christopher D Manning. “Glove: Global vectors for word representation.” In: *Proceedings of the 2014 Conference on Empirical Methods in Natural Language Processing (EMNLP)*. 2014, pp. 1532–1543.
- [15] Jacob Devlin et al. “BERT: Pre-training of deep bidirectional transformers for language understanding.” In: *arXiv preprint arXiv:1810.04805* (2018).
- [16] Sören Auer et al. “Dbpedia: A nucleus for a web of open data.” In: *The semantic web* (2007), pp. 722–735.
- [17] Denny Vrandečić and Markus Krötzsch. “Wikidata: a free collaborative knowledge base.” In: *Communications of the ACM* 57.10 (2014), pp. 78–85.
- [18] Quan Wang et al. “Knowledge graph embedding: A survey of approaches and applications.” In: *IEEE Transactions on Knowledge and Data Engineering* 29.12 (2017), pp. 2724–2743.
- [19] Antoine Bordes et al. “Translating embeddings for modeling multi-relational data.” In: *Advances in neural information processing systems*. Vol. 26. 2013.
- [20] Bishan Yang et al. “Embedding entities and relations for learning and inference in knowledge bases.” In: *International Conference on Learning Representations*. 2015.
- [21] Théo Trouillon et al. “Complex embeddings for simple link prediction.” In: *International Conference on Machine Learning*. 2016, pp. 2071–2080.
- [22] Omar Alonso, Daniel E Rose, and Benjamin Stewart. “Crowdsourcing for relevance evaluation.” In: *ACM SIGIR Forum* 42.2 (2008), pp. 9–15.
- [23] Patricia L Whetzel et al. “BioPortal: enhanced functionality via new Web services from the National Center for Biomedical Ontology to access and use ontologies in software applications.” In: *Nucleic acids research* 39.suppl_2 (2011), W541–W545.
- [24] Egon L Willighagen et al. “The ChEMBL database in 2013: An update.” In: *Nucleic acids research* 41.D1 (2013), pp. D1100–D1107.
- [25] *Earth Science Knowledge Graph*. https://wiki.esipfed.org/Earth_Science_Knowledge_Graph. Accessed: 2025-08-27.
- [26] Sören Auer et al. “GeoKnow: Making the Web an Exploratory Place for Geospatial Data.” In: *Proceedings of the 13th International Semantic Web Conference*. 2014.

- [27] Bernhard Haslhofer et al. “Europeana: Towards a Semantic Search for the European Digital Library.” In: *Library Hi Tech* 31.4 (2013), pp. 584–593.
- [28] Leif Isaksen et al. “Pelagios: A gazetteer and reference system for linking historical place references.” In: *Digital Scholarship in the Humanities* 29.3 (2014), pp. 339–357.
- [29] Amit Singhal. “Introducing the Knowledge Graph: things, not strings.” In: *Official Google Blog* (2012).
- [30] Xin Luna Dong and Divesh Srivastava. *Data integration: The Relational Logic Approach*. Morgan & Claypool Publishers, 2018.
- [31] Zachary C Lipton. “The mythos of model interpretability.” In: *Communications of the ACM* 61.10 (2018), pp. 36–43.
- [32] Alejandro Barredo Arrieta et al. “Explainable Artificial Intelligence (XAI): Concepts, taxonomies, opportunities and challenges toward responsible AI.” In: *Information Fusion* 58 (2020), pp. 82–115.
- [33] Sinno Jialin Pan and Qiang Yang. “A survey on transfer learning.” In: *IEEE Transactions on Knowledge and Data Engineering* 22.10 (2009), pp. 1345–1359.
- [34] Jing Huang et al. “Few-shot learning: A survey.” In: *ACM Computing Surveys* 53.3 (2020), pp. 1–34.
- [35] Paul Voigt and Axel Von dem Bussche. *The EU General Data Protection Regulation (GDPR)*. Springer, 2017.
- [36] Natalya F Noy. “Semantic integration: a survey of ontology-based approaches.” In: *ACM Sigmod Record* 33.4 (2004), pp. 65–70.
- [37] Pavel Shvaiko and Jérôme Euzenat. “Ontology matching: state of the art and future challenges.” In: *IEEE Transactions on Knowledge and Data Engineering* 25.1 (2013), pp. 158–176.
- [38] Thomas Homburg. “INSPIRE: Infrastructure for Spatial Information in the European Community.” In: *Geospatial World Forum*. 2012.